


# Operating Systems

## Programming Assignment #5

### Page Replacement Simulation: LRU and CFLRU

Prof. Li-Pin Chang  
CS@NYCU

# Simulation



# Trace File Format

[Op] [Byte-offset]


- Separated by a space
- Op is either “R” or “W”
- 64-bit byte offset in hex
- Convert the byte offset into page number
  - Cache unit = page size = 4KB
- Total references = 49228943

Memory traces of InceptionV3,  
an DNN running on TensorFlow  
for image recognition


```
os_hw5 > inceptionV3_tf_oshw5.txt
 1   R 4020280
 2   R 1fff0004c0
 3   R 4021000
 4   R 1fff000480
 5   R 4021040
 6   R 1fff000440
 7   R 403ae00
 8   R 403aac0
 9   R 4039a80
10   R 4039e80
11   R 4021080
12   R 403ab00
13   R 40210c0
14   R 403ab80
15   R 403ab40
16   R 4021100
17   R 40215c0
18   R 403ad80
19   R 4039ec0
20   R 4039f00
21   R 403abc0
22   R 4039f40
23   R 403ac40
24   R 4039f80
25   R 403acc0
```

# Page Replacement(LRU)

- Example: Frame #=2




# Reference Design for LRU



# Clean-First LRU (CFLRU)


- Like the enhanced second chance algorithm, CFLRU favors clean pages to reduce I/O-write overhead
- Dividing the LRU list into a working region and a clean-first region
- If the working region is full, migrate an LRU page (P4 here) to the clean-first region
- On replacement, from the clean-first region
  1. Select the LRU clean page (P7 here)
  2. If no clean page, select the LRU (dirty) page



# Clean-First LRU (CFLRU)

- Based on your LRU implementation
  - Set the clean-first region size to  $\frac{1}{4}$  of total size
  - E.g., cache size=4096 pages → working region size = 3072 pages
- Details
  - The write back count ++ only when a dirty page is evicted
  - A write-generated new page is a dirty one
- Two design issues
  1. How do you control the clean-first region size?
  2. How do you efficiently find a clean page in the clean-first region? (don't do linear search)

# A Reference Design



# Page Cache Operations

- Page lookup in LRU/CFLRU
- Find a victim in the clean-first region for CFLRU
- Do not use linear search!!!
  - You will receive a grade penalty if you do
  - Implement your own search, or reuse any existing libraries/classes for searching
  - Duplication in this part does not count
  - Note: do not use STL map/unordered map, too slow

# Procedure

1. Read the trace file and load parsed information into a memory buffer
2. Algorithm=LRU
3. For (frame # = 4096, 8192, 16384, 32768 and 65536) do
  - Run simulation
  - Print out page hit count, miss count, page fault ratio, and page write count
4. Print out the total elapsed time of Step 3
  
5. Algorithm=CFLRU
6. For (frame # = 4096, 8192, 16384, 32768 and 65536) do
  - Run simulation
  - Print out page hit count, miss count, page fault ratio, and page write count
7. Print out the total elapsed time of Step 6

# Output Format

LRU policy: (\n)

| Frame | (\t) | Hit | (\t\t) | Miss |
|-------|------|-----|--------|------|
| 4096  | (\t) | %d  | (\t)   | %d   |
| 8192  | (\t) | %d  | (\t)   | %d   |
| 16384 | (\t) | %d  | (\t)   | %d   |
| 32768 | (\t) | %d  | (\t)   | %d   |
| 65536 | (\t) | %d  | (\t)   | %d   |

| (\t\t) | Page fault ratio | (\t)   | Write back count(\n) |
|--------|------------------|--------|----------------------|
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |

Total elapsed time %.6f sec(\n) (\n)

CFLRU policy: (\n)

| Frame | (\t) | Hit | (\t\t) | Miss |
|-------|------|-----|--------|------|
| 4096  | (\t) | %d  | (\t)   | %d   |
| 8192  | (\t) | %d  | (\t)   | %d   |
| 16384 | (\t) | %d  | (\t)   | %d   |
| 32768 | (\t) | %d  | (\t)   | %d   |
| 65536 | (\t) | %d  | (\t)   | %d   |

| (\t\t) | Page fault ratio | (\t)   | Write back count(\n) |
|--------|------------------|--------|----------------------|
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |
| (\t\t) | %.10f            | (\t\t) | %d                   |

Total elapsed time %.6f sec(\n) (\n)

# Output Example

Here, gcc use -o0

```
(base) brian@DESKTOP-881B5Q0:/mnt/c/Users/User/Desktop/碩一workspace/os_hw5$ ./demo_inceptionV3_tf_oshw5.txt
LRU policy:
Frame Hit Miss Page fault ratio Write back count
4096 45406037 3822906 0.0776556588 1598046
8192 47048736 2180207 0.0442870975 868138
16384 48242502 986441 0.0200378261 457798
32768 48844019 384924 0.0078190588 182950
65536 48990300 238643 0.0048476158 105806
Elapsed time: 12.248408 sec

CFLRU policy:
Frame Hit Miss Page fault ratio Write back count
4096 45532699 3696244 0.0750827415 1357444
8192 47007736 2221207 0.0451199409 719727
16384 48233622 995321 0.0202182078 389906
32768 48824097 404846 0.0082237394 167509
65536 49005852 223091 0.0045317040 83062
Elapsed time: 25.687934 sec
```

Your results must be exactly the same as shown here (except time)

Hit: requested page is in the cache; miss: otherwise.

Fault ratio = miss # / total reference #

Write back count = how many dirty pages have been evicted (and requiring write-back)

# Performance Reference

- 12th Gen Intel(R) Core(TM) i7-12700K
  - 12 seconds for LRU
  - 26 seconds for CFLRU
  - InceptionV3
- Use [this](#) to convert your execution time based on the relative CPU performance (single thread)
- If your execution time is way too long, you will receive a score penalty

# Correctness

- Except the elapsed time, your results should be exactly the same as in the output example
- The TAs will prepare **another workload** to validate your implementation
- **LRU 60% CFLRU 40%**
  
- Once again, do not use linear search in anywhere of your program. Your program will definitely run too slow and you will receive a score penalty
- **Tip:** Do not use map/unordered\_map from STL as they are poorly implemented and very slow

# More details

- Total request #:  $\leq 10^8$  references
- Byte offset: 64 bits
- The path+file name of the trace file is an argument of your program (see the screen shot), do not hard-coding the pathname of the trace file
- You can read the trace file into a memory buffer for fast reuse
- Use `gettimeofday()` to get the total elapsed time

# Testing OS Environment

- Ubuntu 22.04+
- Physical installation, VM, or WSL

# Credits

- 呂柏勳 許登豪 helped design this project
- Direct all questions to the current TAs