

Programming Assignment #6

A User-Space File System

Introduction to Operating Systems

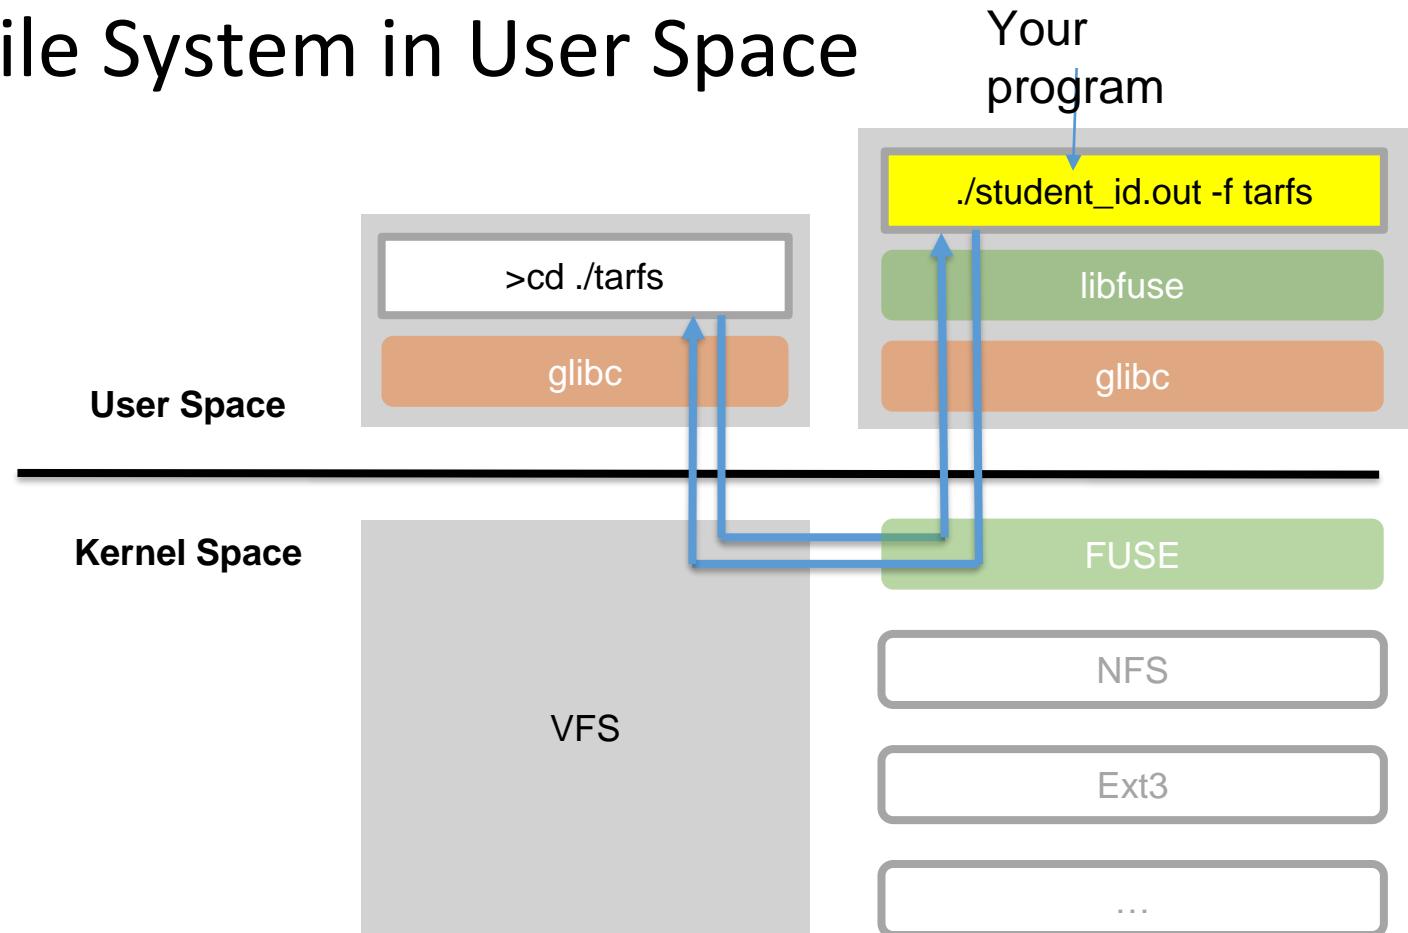
Prof. Li-Pin Chang

CS@NYCU

Objective

- Implementing a user-space file system that mounts a tar file to a specified directory
 - Files in the tar files can be accessed through the directory tree
- This assignment is based on FUSE of Linux
 - Your program will run as a FUSE server
 - Test your FUSE server from another terminal

FUSE – File System in User Space



- FUSE: a kernel component plus a user-space library
- Purpose: accessing existing files/services through the file system interface
 - E.g., an FTP file system, a zip file system, etc.

The Complete FUSE Operation Set

```
int(* getattr )(const char *, struct stat *, struct fuse_file_info *fi)
int(* readlink )(const char *, char *, size_t)
int(* mknod )(const char *, mode_t, dev_t)
int(* mkdir )(const char *, mode_t)
int(* unlink )(const char *)
int(* rmdir )(const char *)
int(* symlink )(const char *, const char *)
int(* rename )(const char *, const char *, unsigned int flags)
int(* link )(const char *, const char *)
int(* chmod )(const char *, mode_t, struct fuse_file_info *fi)
int(* chown )(const char *, uid_t, gid_t, struct fuse_file_info *fi)
int(* truncate )(const char *, off_t, struct fuse_file_info *fi)
int(* open )(const char *, struct fuse_file_info *)
int(* read )(const char *, char *, size_t, off_t, struct fuse_file_info *)
int(* write )(const char *, const char *, size_t, off_t, struct fuse_file_info *)
int(* statfs )(const char *, struct statvfs *)
int(* flush )(const char *, struct fuse_file_info *)
int(* release )(const char *, struct fuse_file_info *)
int(* fsync )(const char *, int, struct fuse_file_info *)
int(* setxattr )(const char *, const char *, const char *, size_t, int)
int(* getxattr )(const char *, const char *, char *, size_t)
int(* listxattr )(const char *, char *, size_t)
int(* removexattr )(const char *, const char *)
int(* opendir )(const char *, struct fuse_file_info *)
int(* readdir )(const char *, void *, fuse_fill_dir_t, off_t, struct fuse_file_info *, enum fuse_readdir_flags)
int(* releasedir )(const char *, struct fuse_file_info *)
int(* fsyncdir )(const char *, int, struct fuse_file_info *)
void *(* init )(struct fuse_conn_info *conn, struct fuse_config *cfg)
void(* destroy )(void *private_data)
int(* access )(const char *, int)
int(* create )(const char *, mode_t, struct fuse_file_info *)
int(* lock )(const char *, struct fuse_file_info *, int cmd, struct flock *)
int(* utimens )(const char *, const struct timespec tv[2], struct fuse_file_info *fi)
int(* bmap )(const char *, size_t blocksize, uint64_t *idx)
int(* ioctl )(const char *, unsigned int cmd, void *arg, struct fuse_file_info *, unsigned int flags, void *data)
int(* poll )(const char *, struct fuse_file_info *, struct fuse_pollhandle *ph, unsigned *revents)
int(* write_buf )(const char *, struct fuse_bufvec *buf, off_t off, struct fuse_file_info *)
int(* read_buf )(const char *, struct fuse_bufvec **bufp, size_t size, off_t off, struct fuse_file_info *)
int(* flock )(const char *, struct fuse_file_info *, int op)
int(* fallocate )(const char *, int, off_t, off_t, struct fuse_file_info *)
ssize_t(* copy_file_range )(const char *path_in, struct fuse_file_info *fi_in, off_t offset_in, const char *path_out, off_t(* lseek )(const char *, off_t off, int whence, struct fuse_file_info *)
```

Specification

- Your server must support the following operations
 - Listing directories
 - Reading files
 - Handling soft link
- Files from us
 - hw6.zip that contains everything: tar files, test script, etc.
- You write
 - A file server “{student_ID}.c” or “{student_ID}.cpp”
- Notice
 - The input filename is hardcoded as “**test.tar**” (you can check the script)
 - You must turn in **one** c/cpp file only

Necessary FUSE Operations

```
struct fuse_operations {  
    int (*readdir)(const char *, void *, fuse_fill_dir_t, off_t, struct fuse_file_info *);  
    int (*getattr)(const char *, struct stat *);  
    int (*read)(const char *, char *, size_t, off_t, struct fuse_file_info *);  
    int (*readlink)(const char *path, void *buffer, size_t size);  
    //many other functions...  
}
```

- The complete FUSE operation set contains many callback functions, but only three are necessary to this assignment
 - **readdir**: Get a list of files and directories that reside in the directory. (Get file names only)
 - **getattr**: Get attributes of a file/directory.
 - **read**: Get the content of a file
 - **readlink**: get a symbolic link
- Leave null to the other operations

readdir

```
int readdir(const char *path, void *buffer, fuse_fill_dir_t filler, off_t offset, struct fuse_file_info *fi);
```

- **Arguments**

- path: (full) relative path to the file/directory.
- buffer: store file names into this buffer using the provided filler
- filler: FUSE callback function to fill file names into the buffer, e.g.,
 - filler(buffer, “file1.txt”, NULL, 0);
 - filler(buffer, “dir1”, NULL, 0);
 - The function will handle internal buffer organization
- offset and fi: Not used in this assignment

- **Return values**

- Always return 0.

getattr

```
int getattr(const char *path, struct stat *st);
```

- **Arguments**
 - path: (full) relative path to the file/directory.
 - st: You should fill the necessary fields of this structure.
- About structure stat: <https://pubs.opengroup.org/onlinepubs/009695399/basedefs/sys/stat.h.html>
- Necessary Fields of st: **st_uid, st_gid, st_mtime, st_size and st_mode**
 - st_mode of the root directory (“/”) should be set to: S_IFDIR | 0444 (act like a read only directory)
 - Other directories: S_IFDIR | accessMode
 - Regular files: S_IFREG | accessMode
- **Return values**
 - Return 0 on success.
 - Return a nonzero value on failure. (If cannot find the specified file/directory)

read

```
int read(const char *path, char *buffer, size_t size, off_t offset, struct fuse_file_info *fi);
```

- **Arguments**

- path: (full) relative path to the file/directory.
- buffer: You should store the requested file content into this buffer.
- size: Max # of chars to store in the buffer. (Should not overrun)
- offset: Skip *offset* chars from the beginning of the file and then start reading.
- fi: Not used in this assignment

- **Return values**

- Return number of bytes read successfully.

readlink

```
int readlink(const char *path, void *buffer, size_t size);
```

- **Arguments**

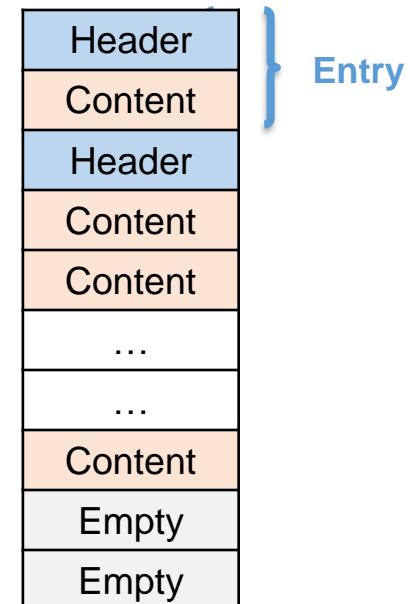
- path: The path of the symbolic link to be resolved
- buffer: A buffer to store the target of the symbolic link (the path it points to)
- size_t: The maximum size of the buffer

- **Return values**

- Always return 0.

Tar File Format

- A tar file contains a series of entries, each of which contains a header and contents
 - One entry per file
 - Header: metadata of a file
 - Contents: contents of the file
- You are responsible for reading and parsing the information of a tar file
- Detailed explanation of tar format:
<https://www.systutorials.com/docs/linux/man/5-tar/>



Skeleton of Your FUSE Server

```
#define FUSE_USE_VERSION 30

#include <fuse.h>
#include <string.h>

int my_readdir(const char *path, void *buffer, fuse_fill_dir_t filler, off_t offset, struct fuse_file_info *fi) { /*do something*/ }
int my_getattr(const char *path, struct stat *st) { /*do something*/ }
int my_read(const char *path, char *buffer, size_t size, off_t offset, struct fuse_file_info *fi) { /*do something*/ }
int readlink(const char *path, void *buffer, size_t size) {/* do something */ }

static struct fuse_operations op;

int main(int argc, char *argv[])
{
    memset(&op, 0, sizeof(op));
    opgetattr = my_getattr;
    op.readdir = my_readdir;
    op.read = my_read;
    op.readlink = my_readlink;
    return fuse_main(argc, argv, &op, NULL);
}
```

Remarks on readlink (soft link)

- The callback ***readlink*** handles symbolic links
- TAR has dedicate records for symbolic link, check the TAR format
- Don't worry about dangling link. The target file is always present in our test cases.
- Your symbolic link should retain its attributes, and the test script must be able to read the target file correctly

Compiling Your FUSE server

- Install FUSE for your Ubuntu VM
 - `sudo apt install libfuse-dev`
- Compile

```
$ gcc student_id.c -o student_id.out `pkg-config fuse --cflags --libs`
```

OR

```
$ g++ student_id.cpp -o student_id.out `pkg-config fuse --cflags --libs`
```

Running and Testing Your FUSE server

這是手動測試的方法，助教評分會用下頁的自動測試腳本！

Terminal 1: Run your fuse server

```
$ ./student_id.out -f tarfs
```

(Ctrl+C to abort)

(you may print some messages here)

Your server reads and parses file
information in test.tar

test.tar

Kernel sends FUSE requests
to your user-space server

Your server responds to FUSE requests
(readdir, read, getarrt) based on the parsed
information from test.tar

Terminal 2: Access your file system

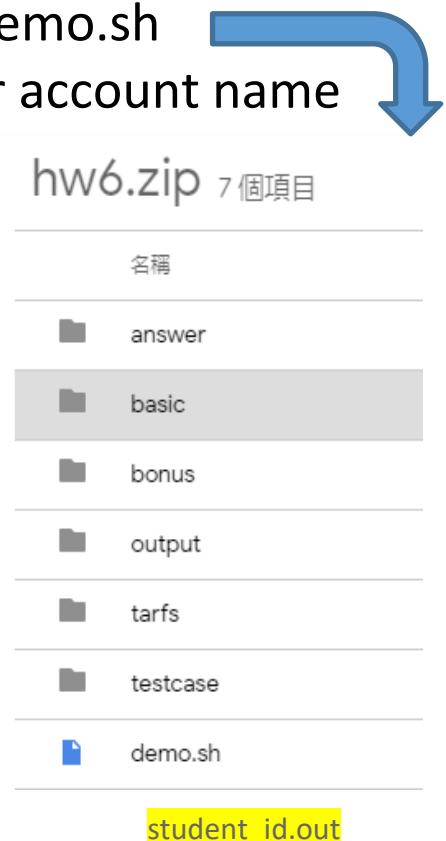
```
$ cd tarfs (get into the mount point)
$ ls (list file names)
file1.c file2.txt dir1 dir2
$ ls -l (list more attributes)
total 0
-rw-r--r-- 0 useraaa useraaa 1234 Oct 10 13:37 file1.c
-rw-r--r-- 0 useraaa useraaa 2244 Oct 10 13:11 file2.txt
drw-r--r-- 0 useraaa useraaa 0 Oct 10 12:55 dir1
drw-r--r-- 0 useraaa useraaa 0 Oct 10 12:11 dir2
$ cat file2.txt (print out file content)
Hello world!
```

“tarfs” is an empty directory, used as a
mount point of your FUSE server

Test Script

1. Download and extract hw6.zip
2. Put your executable file in the same directory as demo.sh
3. In the following files, change “**nctuos**” to your user account name
 - ./answer/1.txt
 - ./answer/2.txt
 - ./answer/5.txt
4. Add an executable attribute to ./ testcase/*txt

```
/home/nctuos/Documents/tarfs
/home/nctuos/Documents/tarfs/dir1/dir2
/home/nctuos/Documents/tarfs
./testcase/1.txt: line 7: cd: largefiles:
/home/nctuos/Documents/tarfs/dir
```



1. Run the script: ./demo.sh <pathname of your FUSE server>

Testing Results

```
› (base) yuiwu@YuiWu-DESKTOP:~/os_assignment/test$ ./demo.sh hw6.out
===== basic case 1 =====
Your answer is correct
===== basic case 2 =====
Your answer is correct
===== basic case 3 =====
Your answer is correct
===== basic case 4 =====
Your answer is correct
===== softlink case 5 =====
Your answer is correct
===== softlink case 6 =====
Your answer is correct
===== Summary =====
[Correct]: 1 2 3 4 5 6
```



```
● (base) yuiwu@YuiWu-DESKTOP:~/os_assignment/test$ ./demo.sh hw6.out
===== basic case 1 =====
Your answer is wrong
===== basic case 2 =====
Your answer is correct
===== basic case 3 =====
Your answer is correct
===== basic case 4 =====
Your answer is correct
===== softlink case 5 =====
Your answer is correct
===== softlink case 6 =====
Your answer is correct
===== Summary =====
[Correct]: 2 3 4 5 6
[ Wrong ]: 1
```

All pass

Some errors

Basic: 1,2,3,4

Softlink: 5,6

Remarks

- If you get a broken mount point during testing, use the following command to force unmount
 - `sudo umount -l <your_mount_point>`
- Do not use external library to parse tar files; parse on your own!
- Do not untar files from `test.tar` and copy them to the mount point... this is cheating!!!

Testing OS Environment

- Ubuntu 22.04+
- Physical installation, VM, or WSL

Credits

- 吳雅柔 吳宥毅 helped design this assignment
- Questions should be directed to the current TAs